

Part 5

System Calls ,Services and Processes

- A **system call** is the programmatic way in which a computer program requests **a service** from the kernel of the operating system it is executed on.
- A system call is a way for programs to **interact with the operating system**.
- A System call **provides** the services of the operating system to the user programs via **Application Program Interface(API)**.
- It provides an interface between a process and operating system to allow user-level processes to request services of the operating system.
- System calls are the only **entry points** into the kernel system.
- All programs needing resources must use **system calls**.

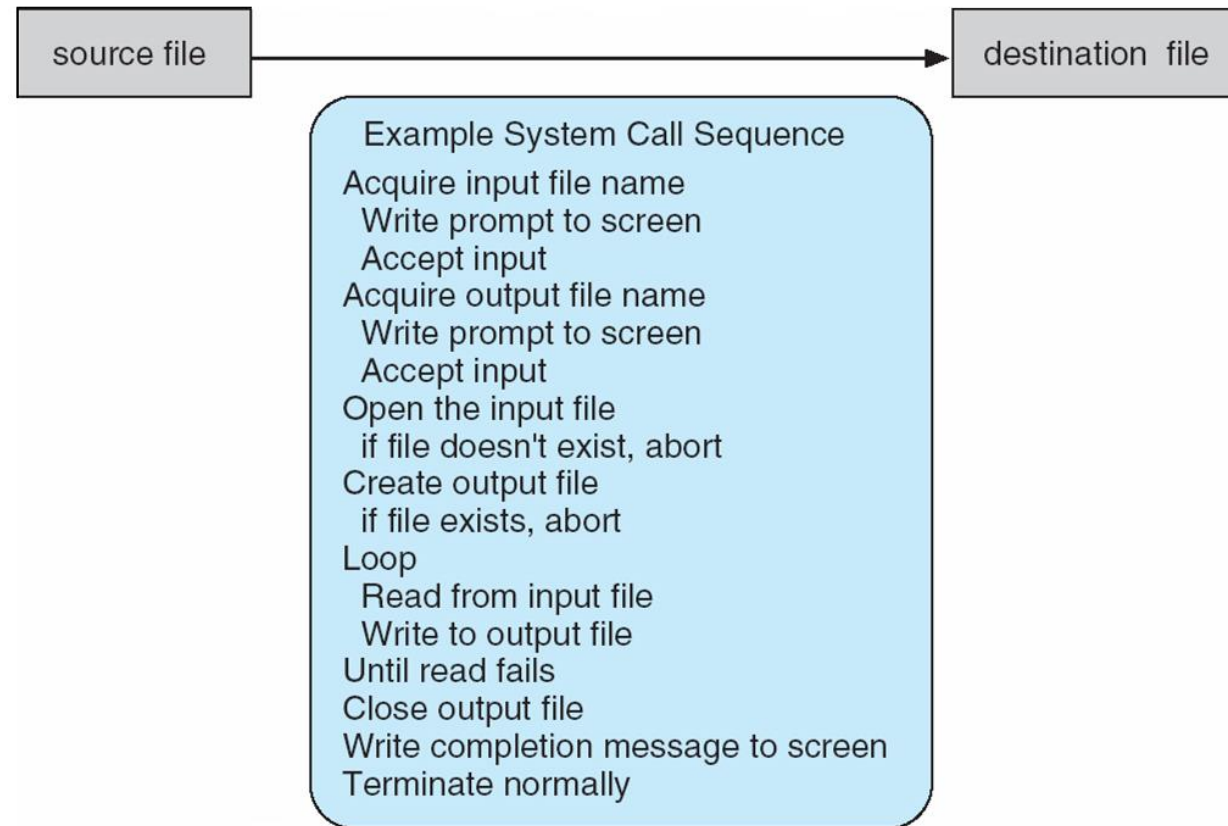
System Calls

- **Programming interface to the services provided by the OS**
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use
- Three most common **APIs** are **Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)**

Note that the system-call names used throughout this text are generic

Example of System Calls

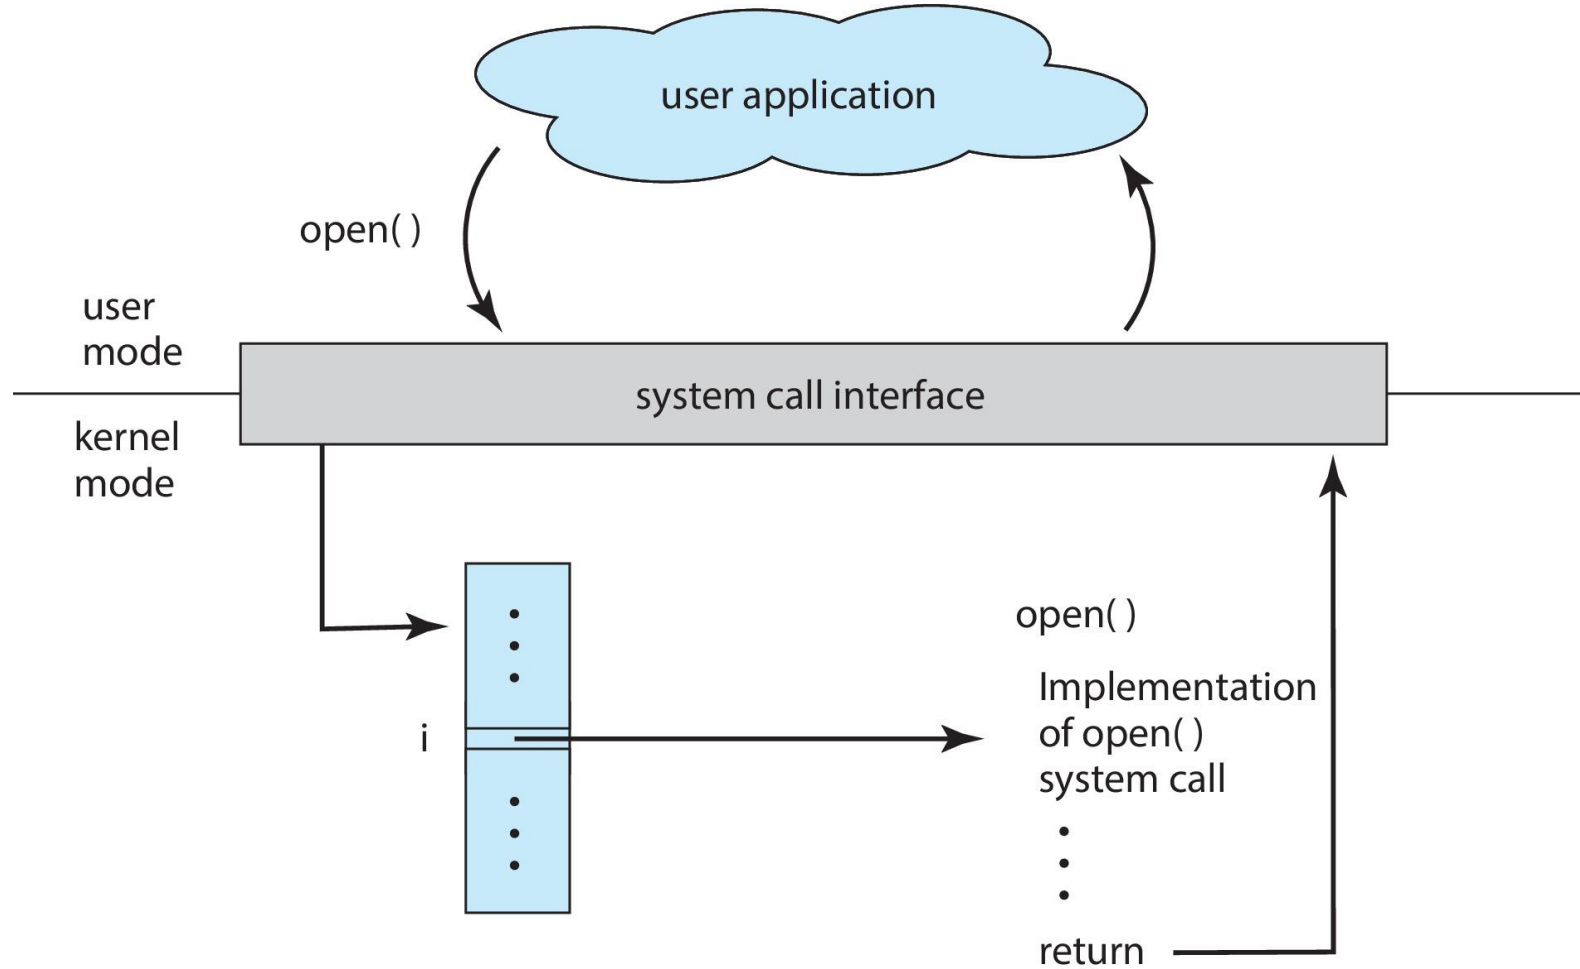
- System call sequence to copy the contents of one file to another file



System Call Implementation

- Typically, **a number associated with each system call**
 - **System-call interface** maintains a table indexed according to these numbers
- The system call interface invokes the intended system call in OS kernel and **returns status of the system call and any return values**
- The caller need know nothing about how the system call is implemented
 - Just needs to **obey** API and understand what OS will do as a result call
 - Most details of OS interface hidden from programmer by API
 - ▶ Managed by run-time support library (set of functions built into libraries included with compiler)

API – System Call – OS Relationship



Types of System Calls

□ **Process control**

- create process, terminate process
- end, abort
- load, execute
- get process attributes, set process attributes
- wait for time
- wait event, signal event
- allocate and free memory
- Dump memory if error
- **Debugger** for determining **bugs, single step** execution
- **Locks** for managing access to shared data between processes

Types of System Calls (cont.)

□ File management

- create file, delete file
- open, close file
- read, write, reposition
- get and set file attributes

□ Device management

- request device, release device
- read, write, reposition
- get device attributes, set device attributes
- logically attach or detach devices

Types of System Calls (Cont.)

□ Information maintenance

- get time or date, set time or date
- get system data, set system data
- get and set process, file, or device attributes

□ Communications

- create, delete communication connection
- send, receive messages if **message passing model** to **host name** or **process name**
 - ▶ From **client** to **server**
- **Shared-memory model** create and gain access to memory regions
- transfer status information
- attach and detach remote devices

Types of System Calls (Cont.)

□ Protection

- Control access to resources
- Get and set permissions
- Allow and deny user access

Examples of Windows and Unix System Calls

EXAMPLES OF WINDOWS AND UNIX SYSTEM CALLS

The following illustrates various equivalent system calls for Windows and UNIX operating systems.

	Windows	Unix
Process control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File management	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device management	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communications	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



statcounter

GlobalStats



Win10

67.46%

Win11

26.52%

Win7

3.34%

Win8.1

1.66%

WinXP

0.64%

Win8

0.28%

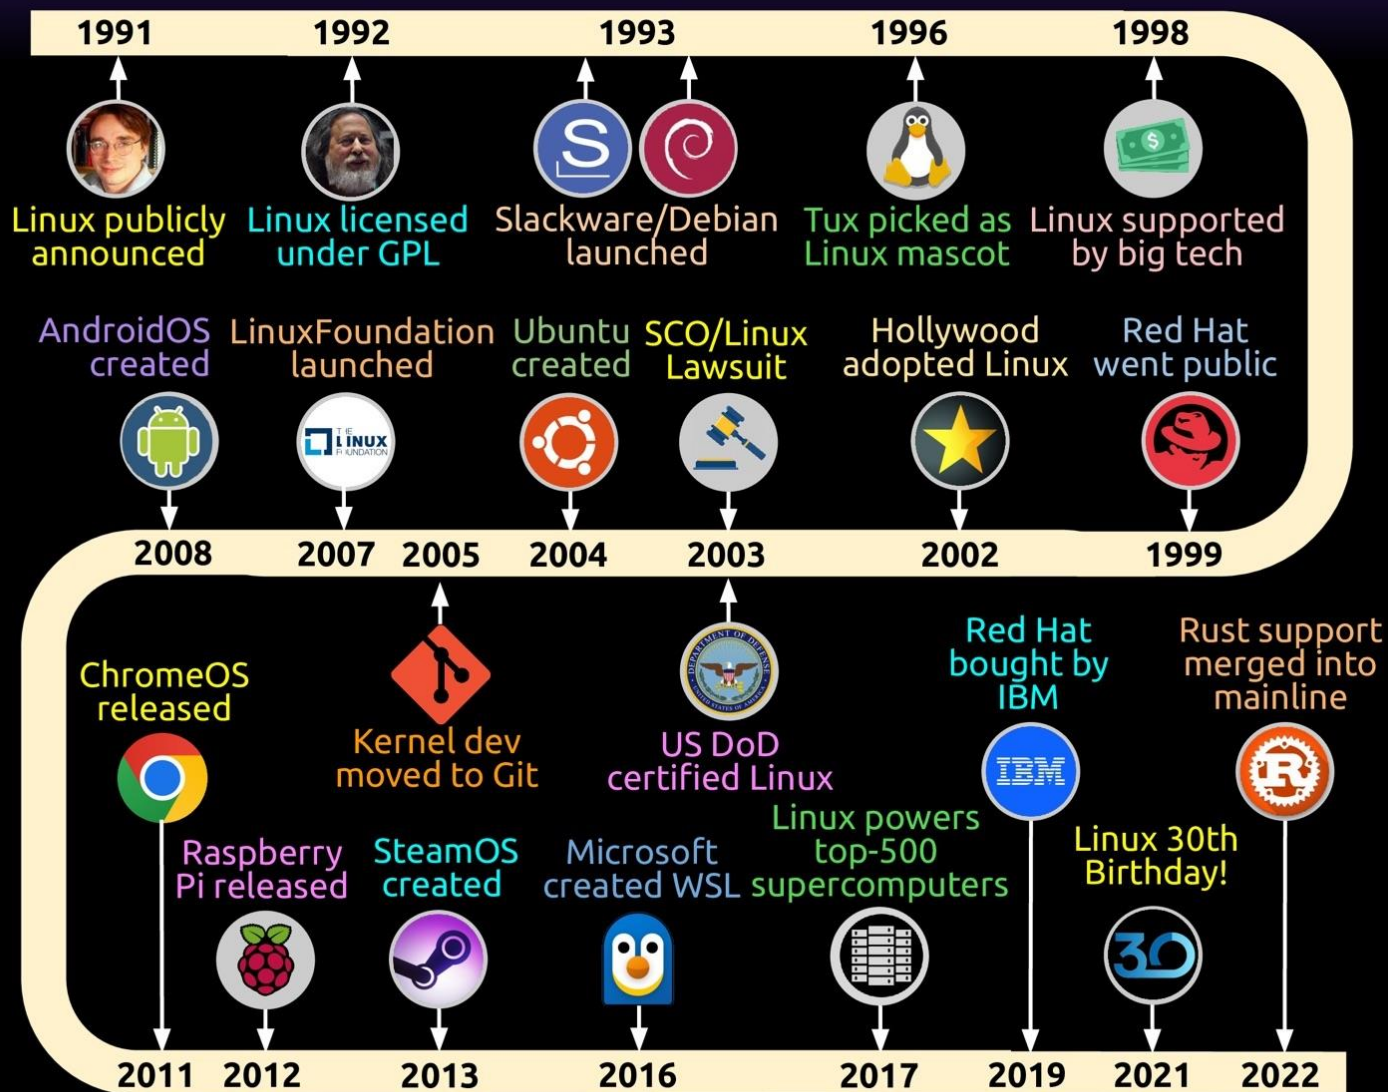
**Desktop Windows Version Market Share
Worldwide - December 2023**



Memorable Events in Linux History

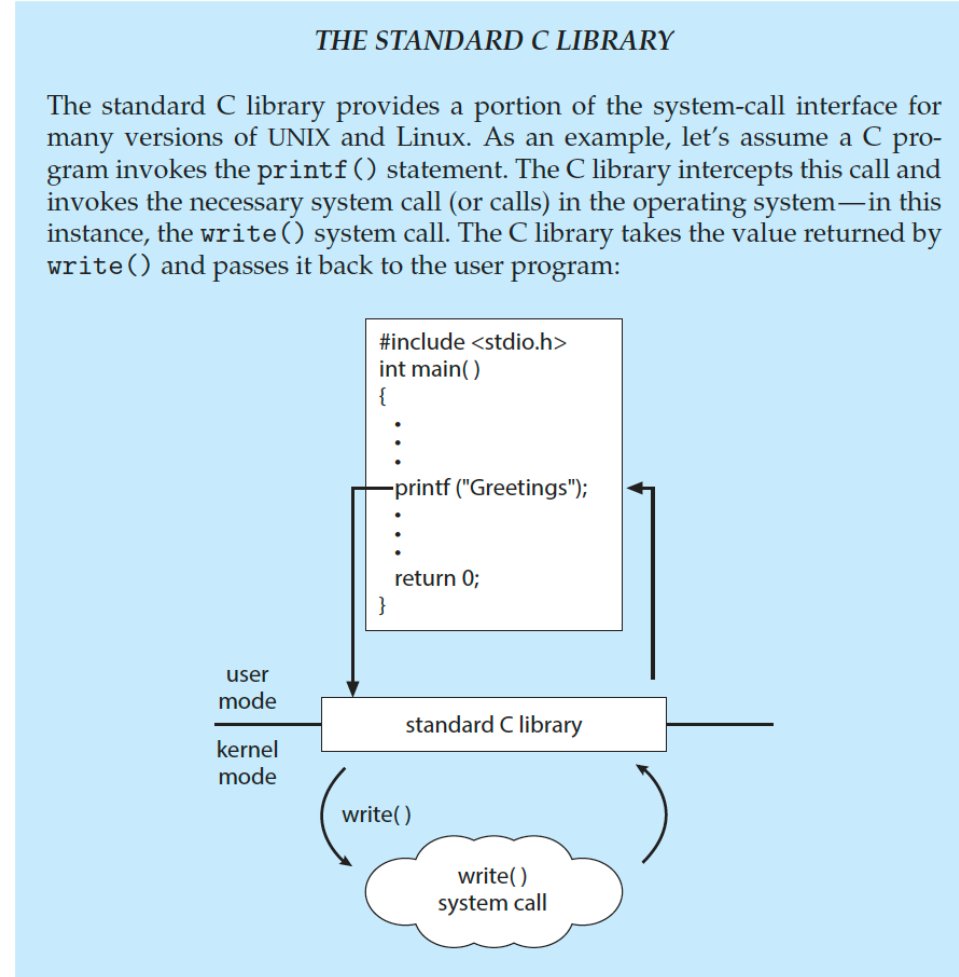


Created by [@dan_nanni](#) on Instagram



Standard C Library Example

- C program invoking `printf()` library call, which calls `write()` system call



System Services

- System programs provide a convenient environment for program **development and execution**. They can be divided into:
 - File manipulation
 - Status information sometimes stored in a file
 - Programming language support
 - Program loading and execution
 - Communications
 - Background services
 - Application programs
- **Most users' view of the operation** system is defined by system programs, not the actual system calls

System Services (cont.)

- Provide a convenient environment for program development and execution
 - Some of them are simply user interfaces to system calls; others are considerably more complex
- **File management** - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- **Status information**
 - Some ask the system for info - date, time, amount of available memory, disk space, number of users
 - Others provide detailed performance, logging, and debugging information
 - Typically, these programs format and print the output to the terminal or other output devices
 - Some systems implement a **registry** - used to store and retrieve configuration information

System Services (Cont.)

- ❑ **File modification**
 - ❑ Text editors to create and modify files
 - ❑ Special commands to search contents of files or perform transformations of the text
- ❑ **Programming-language support** - Compilers, assemblers, debuggers and interpreters sometimes provided
- ❑ **Program loading and execution**- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language
- ❑ **Communications** - Provide the mechanism for creating virtual connections among processes, users, and computer systems
 - ❑ Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another

System Services (Cont.)

□ Background Services

- Launch at boot time
 - ▶ Some for system startup, then terminate
 - ▶ Some from system boot to shutdown
- Provide facilities like disk checking, process scheduling, error logging, printing
- Run in user context not kernel context
- Known as **services**, **subsystems**, **daemons**

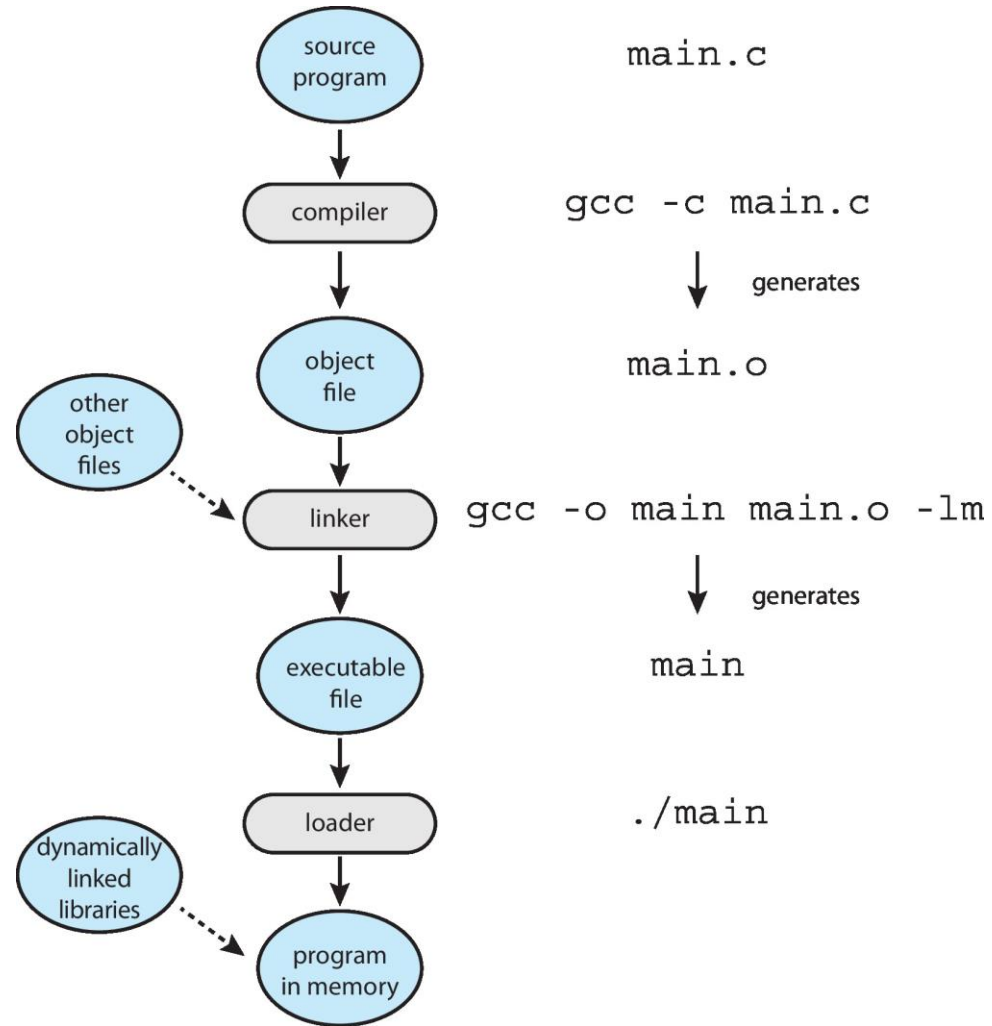
□ Application programs

- Don't pertain to system
- Run by users
- Not typically considered part of OS
- Launched by command line, mouse click, finger poke

Linkers and Loaders

- ❑ Source code compiled into object files designed to be loaded into any physical memory location – **relocatable object file**
- ❑ **Linker** combines these into single binary **executable** file
 - ❑ Also brings in libraries
- ❑ Program resides on secondary storage as binary executable
- ❑ Must be brought into memory by **loader** to be executed
 - ❑ **Relocation** assigns final addresses to program parts and adjusts code and data in program to match those addresses
- ❑ Modern general purpose systems don't link libraries into executables
 - ❑ Rather, **dynamically linked libraries** (in Windows, **DLLs**) are loaded as needed, shared by all that use the same version of that same library (loaded once)
- ❑ Object, executable files have standard formats, so operating system knows how to load and start them

The Role of the Linker and Loader



Why Applications are Operating System Specific

- ❑ **Apps compiled on one system usually not executable on other operating systems**
- ❑ Each operating system provides its own unique system calls
 - ❑ Own file formats, etc
- ❑ **Apps can be multi-operating system**
 - ❑ Written in interpreted language like **Python, Ruby**, and **interpreter** available on multiple operating systems
 - ❑ App written in language that includes a VM containing the running app (like Java)
 - ❑ Use standard language (like C), compile separately on each operating system to run on each
- ❑ **Application Binary Interface (ABI)** is architecture equivalent of API, defines how different components of binary code can interface for a given operating system on a given architecture, CPU, etc

Operating System Design and Implementation

- Design and Implementation of OS not “solvable”, but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- **Start the design by defining goals and specifications**
- Affected by choice of hardware, type of system
- **User** goals and **System** goals
 - **User goals** – operating system should be convenient to use, easy to learn, reliable, safe, and fast
 - **System goals** – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient

Operating System Design and Implementation (Cont.)

- Important principle to separate
Policy: *What* will be done?
Mechanism: *How* to do it?
- **Mechanisms determine how to do something, policies decide what will be done**
- The separation of policy from mechanism is a very important principle, it allows **maximum flexibility** if policy decisions are to be changed later (example – timer)
- **Specifying and designing an OS is highly creative task of software engineering**

Operating System Structure

- General-purpose OS is very large program
- Various ways to structure ones
 - Simple structure – MS-DOS
 - More complex -- UNIX
 - Layered – an abstraction
 - Microkernel -Mach

MAJOR DESIGN ISSUES

OVERVIEW

- Transparency
- Flexibility
- Reliability
- Performance
- Scalability
- Naming
- Replication
- Synchronization
- Communication Model
- Security

Transparency

Multiple Computers are used but the user gets a view of only single system being used. Makes the network invisible to user/applications.

Various degrees of transparency

- ❑ Access Transparency
- ❑ Location Transparency
- ❑ Name Transparency
- ❑ Data Transparency
- ❑ Execution Transparency
- ❑ Performance Transparency

Flexibility

- Flexible operating systems are taken to be those whose designs have been motivated to some degree by the desire to allow the system to be tailored, either statically or dynamically, to the requirements of specific applications or application domains.
- The use of **object orientation** is a common feature of many flexible operating systems.

Reliability

- ❑ In general, **reliability** is the ability of a person or system to perform and maintain its functions in routine circumstances, as well as hostile or unexpected circumstances.
- ❑ Reliability is generally considered **important by end users**.
- ❑ Not all companies making operating systems have a similar standard. Even among operating systems where reliability is a priority, there is a range of quality.
- ❑ Also, an operating system may be extremely reliable at one kind of task and extremely unreliable at another.

- Current operating systems have two characteristics that make them unreliable and insecure.

- 1> **They are huge**

- 2> **They have very poor fault isolation.**

- The Linux kernel has more than 2.5 million lines of code; the Windows XP kernel is more than twice as large.

Performance

- The performance of computer hardware typically increases **monotonically with time**.
- Even if the same *could* be said of software, the rate at which software performance improves is usually very slow compared to that of hardware.
- In fact, many might opine that there is plenty of software whose performance has *deteriorated* consistently with time.
- Moreover, it is rather difficult to establish an objective performance metric for software as complex as an operating system: a "*faster OS*" is a *very* subjective, context dependent phrase.
- How to increase performance of an Operating System ?

- ❑ **Apple** computers **do not hibernate**. Rather, when they "sleep", enough devices (in particular, the dynamic RAM) are kept alive (at the cost of some battery life, if the computer is running on battery power).
- ❑ Consequently, upon wakeup, the user perceives instant-on behavior: a very desirable effect.
- ❑ Similarly, by default the system tries to keep network connections alive even if the machine sleeps.
- ❑ For example, if you login (via SSH, say) from one PowerBook to another, and both of them go to sleep, your login should stay alive within the constraints of the protocols.

Scalability

- **Scalability** is a desirable property of a system, a network, or a process, which indicates its ability to either handle growing amounts of work in a graceful manner, or to be readily enlarged.
- A system whose performance improves after adding hardware, proportionally to the capacity added, is said to be a **scalable system**.

Naming

- The resources in a distributed system are spread across different computers and a naming scheme has to be devised so that users can discover and refer to the resources that they need.
- An example of such a naming scheme is the URL (Uniform Resource Locator) that is used to identify WWW pages. If a meaningful and universally understood identification scheme is not used then many of these resources will be inaccessible to system users.

Replication

- Replication is one of the oldest and most important topics in the overall area of distributed systems.
- Whether one replicates data or computation, the objective is to have some group of processes that handle incoming events.
- If we replicate data, these processes are passive and operate only to maintain the stored data, reply to read requests, and apply updates.
- When we replicate computation, the usual goal is to provide fault-tolerance.

Synchronization

- Processes require coordination to achieve synchronization

- Types of synchronization:

- Barrier synchronization

- Process must reach a common synchronization point before they can continue.

- Condition coordination

- A process must wait for a condition that will be set asynchronously by other interacting processes to maintain some ordering of execution.

- Mutual exclusion

- Concurrent processes must have mutual exclusion when accessing a critical shared resource.

Communication Model

- Lower level: message passing
- Higher level logical communication provides transparency.
 - Client/server model communication.
 - ▶ All system communication are seen as a pair of message exchanges between the client and server.
 - Remote Procedure Call, (RPC), communication.
 - ▶ RPC built on top of client/server model.
- Request/reply message passing as used in programming procedure-call concept.

SECURITY

□ Authentication

Clients , Servers and messages must be authenticated.

□ Authorization

Access control has to be performed across a physical network with heterogeneous components under different administrative units using different security models.

- ❑ Operating system authentication (OS authentication) is a way of using operating system login credentials to authenticate database users.
- ❑ One aspect of OS authentication can be used to authenticate database administrators.
- ❑ OS authentication is needed because there must be a way to identify administrative users even if the database is shut down.

Types of Operating System:

- Normal Operating System

An operating system, or OS, is a software program that enables the computer hardware to communicate and operate with the computer software.

- Distributed Operating Systems

Integration of system services presenting a transparent view of a multiple computer system with distributed resources and control. Consisting of concurrent processes accessing distributed shared or replicated resources through message passing in a network environment.

□ Multiprocessor Operating Systems

For the most part, multiprocessor operating systems are just regular operating systems. They handle system calls, do memory management, provide a file system, and manage I/O devices. Nevertheless, there are some areas in which they have unique features. These include process synchronization, resource management, and scheduling.

□ Database Operating Systems

Database systems place increased demands on an operating system to efficiently support: concept of a transactions, manage large volumes of data, concurrency control, system failure control.

□ Real-time Operating Systems

A real-time operating system is a multitasking operating system intended for real-time applications.

Such applications include embedded systems (programmable thermostats, household appliance controllers, mobile telephones), industrial robots, spacecraft, industrial control.

The Big Pic of OS

